

R jobs in the heavens

Running R jobs at large scale (remotely)

- A good point: Collaboration and sharing codes and analysis
- Local clusters... you may have access to one in you institution, but
 - high cost for maintenance and administration...
 - it must be accessible to all researchers... independently from subject
 - young guns and basic sciences are not very appealing to deserve access...

Cloud possibilities

Amazon Web Services - AWS

- Getting started with R on Amazon Web Services - <https://aws.amazon.com/blogs/opensource/getting-started-with-r-on-amazon-web-services/>
- Running an R Code on AWS Batch - <https://medium.com/geekculture/running-an-r-container-on-aws-batch-on-production-9be336c34f95>
 - Jobs are submitted via Docker files with help of **packrat** R package
 - cost? only free trial?

Azure - Microsoft

- R workloads on Azure Batch - <https://azure.microsoft.com/pt-br/blog/r-workloads-on-azure-batch/>
- Tutorial: Run a parallel workload with Azure Batch using the .NET API - <https://docs.microsoft.com/en-us/azure/batch/tutorial-parallel-dotnet>
 - paid if you have an institutional account?

Google Cloud Platform

- Running R at Scale on Compute Engine - clusters setup - <https://cloud.google.com/architecture/running-r-at-scale>
 - If not eligible for a free trial, "... a 6-node cluster composed of n1-standard-4 instances, would cost \$1.84/hr ..."

Oracle cloud computing

- Using Oracle R Enterprise Embedded R Execution
 - free credits for new accounts... better performance? lower costs?

R possibilities

R Cloud - AT&T Labs

- Try It, Online or Locally - <https://rcloud.social/tryit/index.html>
 - support to different programming languages in the same code
 - code and analysis everything is public

Rstudio Server or connect???

- From “What is the difference between rstudio-connect and rstudio-pro” - <https://community.rstudio.com/t/what-is-the-difference-between-rstudio-connect-and-rstudio-pro/43949>
- RStudio Server Pro - Very similar to the Desktop IDE, but runs on a server.
 - it allows you to have more compute resources closer to your data,
 - it provides a uniform environment for teams that makes it easier to collaborate,
 - and it provides controls for administrators to monitor and scale work.
- RStudio Connect - Makes it easy to share R Markdown reports,
 - deploy shiny web applications, and APIs written in R.
 - GitHub stores your code statically, RStudio Connect knows how to run it, (accessibility to non R users... shiny applications)
 - run reports on a schedule and send emails with the results.
- (easy way) Without Docker!!!
 - R Studio Server on Google Cloud - <https://towardsdatascience.com/r-studio-server-on-google-cloud-dd69b8bff80b>
 - Deploying R Studio on Compute Engine - <https://sookocheff.com/post/r/deploying-r-studio-to-compute-engine/>
- (easiest way) pre-configured from Marketplace
 - Ready-to-use and most of the time more expensive
 - Getting Started with RStudio Connect for GCP - <https://support.rstudio.com/hc/en-us/articles/360033988434-Getting-Started-with-RStudio-Connect-for-GCP>

Docker

- All starts with Docker;
- Primary focused on reproducibility? But also have your code running in any cloud service.

Container it! Rstudio with all you need in a docker container and deploy in a virtual machine... you don't need to install everything again and again manually... and have your code “forever”.

- You have to setup and install everything for the first time, start and stop the job whenever you want.
 - Even if you stop the VM if there is files in the storage will be charged by the service.

Everything is Docker!!!

containering R

- Docker + R - Rocker project - <https://www.rocker-project.org/>
- Rstudio...
 - support for Shiny apps

containering your code

- R.project + Github with Docker...
 - R or Rstudio image;
 - * your favorite R packages;
 - * third-party softwares, GMP e Latte Integrale for example.
 - * your codes;
- dockering R
 - Using R via Rocker A Brief Introduction to Docker for R - http://dirk.eddelbuettel.com/papers/chirug_nov2019_rocker.pdf
 - a good quick Docker introduction for R users - <https://colinfay.me/docker-r-reproducibility/>
 - example of calling additional libraries - Running your R script in Docker - <https://www.r-bloggers.com/2019/02/running-your-r-script-in-docker/>
- dockering R studio
 - a nice introduction and well organised setup - <https://www.symbolix.com.au/blog-main/r-docker-hello>
 - example setting additional libraries (but using rstudio server) - <https://davetang.org/muse/2021/04/24/running-rstudio-server-with-docker/>
 - sharing and Running R code using Docker - <https://aboland.ie/Docker.html>

<https://code.markedmondson.me/r-at-scale-on-google-cloud-platform/>

Third-party software

- How to add C library GMP, GNU for arithmetic precision on GCP? (Or latte distro, both `.tar.gz`)
Via Rstudio Connect not working from its own terminal...
 - latte distro - <https://github.com/latte-int/latte-distro>
 - GMP - <https://gmplib.org/#DOWNLOAD>
- There are R packages for both but them need the distro installed.
 - (*“they are linux... trying to run on a windows machine can be very painful”*)
- Doing this interactively in the example below.
 - how to make it all install from Dockerfile image? (with R base or Rstudio image???)

Example Dockerfile to build RStudio + extras

Example RStudio in a Google cloud VM - interactively

1. to create VM on Google cloud (after having an account)
 - a. VM instances > + Create ...
 - b. create a firewall rule - see details <https://towardsdatascience.com/r-studio-server-on-google-cloud-dd69b8b8ff80b>
2. Open a terminal in browser by clicking SSH button... (how to use Cloud Shell???)
 - a. Installing essential applications and third-party softwares

In the case of Latte distro some packages for installation are not installed in GCP

```
sudo apt update
sudo apt install build-essential
sudo apt-get install m4
```

- b. upload Latte distro `.tar.gz` (terminal has menu to do this by clicking)... unpack and configure it!!!
(latte 1.7.5 from github is not working, maybe some option for `./configure`)

to unpack the bundle

```
tar -xvzf latte-integrale-1.7.3b.tar.gz
```

and to configure, in the unpacked folder type

```
./configure
make
```

Now all the files are in `/home/rstudio-user/latte-integrale-1.7.3b/dest`

3. Installing Docker - thanks to <https://tomroth.com.au/gcp-docker/>
"Before running Dockerfile Docker also needs to be installed."

```
sudo apt update
sudo apt install --yes apt-transport-https ca-certificates curl gnupg2 software-properties-common
curl -fsSL https://download.docker.com/linux/debian/gpg | sudo apt-key add -
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/debian $(lsb_release -cs) stable"
sudo apt update
sudo apt install --yes docker-ce
```

4. Building and running rstudio from your Dockerfile image
 - a. Build it - go to the same directory as your Dockerfile and type the command

```
sudo docker build --rm --force-rm -t rstudio/my_simulation .
```

the `-rm -force-rm` options forces to delete the container once its scripts run or you log out. (stops filling up the server with lots of containers doing nothing.)

`sudo docker image list` if you want to see your image added to the list.

- b. Run it (updated - from <https://hub.docker.com/r/rocker/rstudio>)

```
sudo docker run -d -p 8787:8787 -v $(pwd):/home/rstudio -e PASSWORD=yourpasswordhere rocker/rstudio
```

Done! Now open the rstudio in `localhost:8787` (my localhost is 34.125.71.81) and type `source("my_simulation.R")`. Username is `rstudio` and password `yourpasswordhere`.

```
sudo docker stop my_simulation
```

Rstudio or R on google cloud???

- advantage in using rstudio images... maybe community is working more in containerisation... R Docker faster - “My experiment shows R Docker images will build much faster thanks to the new package manager from RStudio” from <https://medium.com/@skyetetra/r-docker-faster-28e13a6d241d>
- COSTS??? What would all this cost??
 - (easiest way) A VM with 1 vCPU n1-standard-1 (3.75GB RAM and 30GB SSD) with Rstudio from Marketplace costs around $5.5\times$ a VM from Compute > Compute Engine > VM Instances > + Create VM instance in my region;
 - So run your Dockerfile!!!

“the easy way” to get started quickly

(only if youre not confident yet to Dockerisation)

- Rstudio on google cloud Marketplace... <https://support.rstudio.com/hc/en-us/articles/115010260627-Getting-Started-with-RStudio-Workbench-RStudio-Server-Pro-Standard-for-GCP>
 - you can use it all interactively, but has to install everytime you setup a VM...
- Another quick way to launch Rstudio on GCP - via R code
 - Launch RStudio Server in the Google Cloud with two lines of R - <https://code.markedmondson.me/launch-rstudio-server-google-cloud-in-two-lines-r/>

Command line interactively

lots of things didnt work... instal GNU softwares OK... but using rstudio/docker it creates other user and the connection doesnt seem easy for a ...

1. install R... and Run your code...
 - is it possible to run different jobs in different terminals... interactively but not in your machine...
 - how to keep running without being logged in? (... using Rstudio??? ead)

```
sudo apt-get install r-base r-base-dev
```

after installing... type R and... like the old times...

to edit files `sudo apt-get install vim`

latest R base on google cloud???

- Instructions on R page... a bit handy... <http://cloud.r-project.org/bin/linux/debian/#debian-buster-stable>
- A good reference to upgrade R to 4.X... <https://www.charlesbordet.com/en/how-to-upgrade-to-R-4-0-0-on-debian/#>

```
## pre-requisite
sudo apt-get install software-properties-common -y
## force latest and install
sudo apt-key adv --keyserver keyserver.ubuntu.com --recv-keys E298A3A825C0D65DFD57CBB651716619E084DAB9
sudo add-apt-repository 'deb https://cloud.r-project.org/bin/linux/ubuntu focal-cran40/'
sudo apt install r-base
```

Or Install Rstudio Server and have Rstudio facilities

To install Rstudio Server (check the latest version) <https://www.rstudio.com/products/rstudio/download-server/debian-ubuntu/>

```
sudo apt-get install gdebi-core
wget https://download2.rstudio.org/server/bionic/amd64/rstudio-server-1.2.5019-amd64.deb
sudo gdebi rstudio-server-1.2.5019-amd64.deb
sudo apt-get install libcurl4-openssl-dev libssl-dev libxml2-dev
```

Add an user (“R Studio only permits access to users of the system”)

```
sudo adduser rstudio
exit
```

to keep rstudio running when signout, add commands ???not working??? (frist time installed older version, tried rsession.conf but installed as rstudio recommends... and worked :-()

```
auth-timeout-minutes=0
auth-stay-signed-in-days=30
```

to file `/etc/rstudio/rserver.conf` and restart rstudio server

```
sudo rstudio-server restart
```